



CVE-2016-6187 - from zero to root

Exploiting a poison nullbyte in the linux kernel

David Milosevic
June 16, 2024

Agenda

- I. Motivation
- II. Intro to CVE-2016-6187
- III. Exploitation
- IV. Demo
- V. Exploit reliability
- VI. Final words

Motivation

Why should we care about a 8-years old vulnerability?

- initially only meant to be my foot in the door to kernel pwn
- no publicly available exploit that drops a root shell. Though there is a PoC for RIP hijack by Vitaly Nikolenko
- demonstrating the use of novel (?) kernel objects that are fun to exploit

What is CVE-2016-6187 about?

- a single nullbyte overflow in the *setprocattr LSM* hook of *AppArmor*
- affecting kernels up to 4.6.4 (including) with *AppArmor* enabled
- to my knowledge, no popular distro was shipped with a vulnerable configuration

DAC

- enforces classic UNIX permission check on system resources based on ownership and permission bits (*UID*, *GID*)
- on resource access, the kernel first checks *DAC*

MAC

- additional security checks based on configured policies
- access decisions based on central policies instead of resource owners
- security labels for subjects (processes, ...) and resources (files, ...)

AppArmor

- implements mandatory access control (*MAC*)
- restricts access to resources and limits process capabilities based on configured profiles
- it is developed against the *LSM* kernel interface
- currently developed by *Canonical Ltd.*
- other *LSMs* are, for example, *SELinux*, *Smack*, *Tomoyo*, ...

AppArmor

- profiles typically stored in */etc/apparmor.d*
- profiles can be generated using special tooling
- *enforced mode*: current profile enforced and unauthorized access blocked
- *complain mode*: unauthorized access **not** blocked but logged (useful for initial profile creation)

AppArmor

The kernel provides the */proc/self/attr* interface for *LSMs* to implement security labels per process. In the context of *AppArmor*:

- */proc/self/attr/current*: currently enforced profile
- */proc/self/attr/exec*: profile inherited by its child processes
- and a few more...

But to us, only */proc/self/attr/current* will be relevant

Intro to CVE-2016-6187

The bug

```
static ssize_t proc_pid_attr_write(struct file * file, const char __user * buf,
                                  size_t count, loff_t * ppos)
{
    struct inode * inode = file_inode(file);
    void *page;
    ssize_t length;
    struct task_struct *task = get_proc_task(inode);

    length = -ESRCH;
    if (!task)
        goto out_no_task;
    if (count > PAGE_SIZE)
        count = PAGE_SIZE;

    /* No partial writes. */
    length = -EINVAL;
    if (*ppos != 0)
        goto out;

    page = memdup_user(buf, count) [1] alloc count bytes
    if (IS_ERR(page)) {
        length = PTR_ERR(page);
        goto out;
    }

    /* Guard against adverse ptrace interaction */
    length = mutex_lock_interruptible(&task->signal->cred_guard_mutex);
    if (length < 0)
        goto out_free;

    length = security_setprocattr(task,
                                  (char*)file->f_path.dentry->d_name.name,
                                  page, count) [2] call apparmor_setprocattr
```

```
static int apparmor_setprocattr(struct task_struct *task, char *name,
                                void *value, size_t size) [3] size = count
                                value = page
{
    struct common_audit_data sa;
    struct apparmor_audit_data aad = {0,};
    char *command, *args = value; [4] args = value = page
    size_t arg_size;
    int error;

    if (size == 0)
        return -EINVAL;
    /* args points to a PAGE_SIZE buffer, AppArmor requires that
     * the buffer must be null terminated or have size ≤ PAGE_SIZE - 1
     * so that AppArmor can null terminate them
     */
    if (args[size - 1] != '\0') {
        if (size == PAGE_SIZE)
            return -EINVAL;
        args[size] = '\0'; [5] poison nullbyte
    }
}
```

How do we trigger the vulnerability?

```
static const struct file_operations proc_pid_attr_operations = {
    .read      = proc_pid_attr_read,
    .write     = proc_pid_attr_write,
    .llseek    = generic_file_llseek,
};

static const struct pid_entry attr_dir_stuff[] = {
    REG("current",    S_IRUGO|S_IWUGO, proc_pid_attr_operations),
    REG("prev",        S_IRUGO,        proc_pid_attr_operations),
    REG("exec",        S_IRUGO|S_IWUGO, proc_pid_attr_operations),
    REG("fscreate",    S_IRUGO|S_IWUGO, proc_pid_attr_operations),
    REG("keycreate",   S_IRUGO|S_IWUGO, proc_pid_attr_operations),
    REG("sockcreate",  S_IRUGO|S_IWUGO, proc_pid_attr_operations),
};
```

- *proc_pid_attr_write* implements the fops write hook
- trigger the vulnerability by writing to one of the files in */proc/self/attr/*
- for example, */proc/self/attr/current*

How do we trigger the vulnerability?

```
static int apparmor_setprocattr(struct task_struct *task, char *name,
                               void *value, size_t size)
{
    struct common_audit_data sa;
    struct apparmor_audit_data aad = {0,};
    char *command, *args = value;
    size_t arg_size;
    int error;

    if (size == 0)
        return -EINVAL;
    /* args points to a PAGE_SIZE buffer, AppArmor requires that
     * the buffer must be null terminated or have size ≤ PAGE_SIZE - 1
     * so that AppArmor can null terminate them
     */
    if (args[size - 1] != '\0') { [1]
        if (size == PAGE_SIZE) [2]
            return -EINVAL;
        args[size] = '\0';
    }
}
```

- user buffer *args* not null terminated [1]
- total number of bytes written should be less than *PAGE_SIZE* (in our case 4096) [2]

Exploitation

What do we have?

- a heap based poison nullbyte in one of the caches, up to *kmalloc-2048*
- contents of vulnerable buffer fully user controlled

⇒ we can null out the *LSB* of the first quadword of the adjacent chunk

Exploring possible strategies

- nulling out the *LSB* of a function pointer?
 - we could redirect the execution to a single gadget/function relative to the function pointer
 - but very unlikely to find something useful
 - idea quickly discarded
- corrupting the *LSB* of the adjacent freelist pointer?
 - we could shift the pointer if *LSB* $\neq 0 \times 00$ and provoke a ***double free*** scenario
 - from there, we could tamper with other kernel objects
 - sounds like a promising idea

Freelist corruption

At a high level, the plan looks like this

1. groom the heap
2. set up exploit sections
3. trigger the vulnerability
4. perform a double free

Freelist corruption

But first, let's consider the caches we could use

- **caches \geq *kmalloc-256***
 - **LSB** of alloc'd chunk addresses always **0x00**
 - overwriting **0x00** with **0x00** \Rightarrow not an option
- ***kmalloc-96***
 - possible **LSBs**: **60, c0, 20, 80, e0, 40, a0, 00**
 - too much variation \Rightarrow lower reliability
- ***kmalloc-128***
 - possible **LSBs**: **80, 00**
 - 50/50 chance - either we shift the pointer by **0x80** or we overwrite **0x00** with **0x00** and nothing happens
 - we will go for this cache

Exploitation

Freelist corruption - Grooming the Heap

```
[ 6.055636] do_msgsnd: ffff88001d9c9180
[ 6.056035] do_msgsnd: ffff88001d9c9800
[ 6.056415] do_msgsnd: ffff88001d9c9780
[ 6.056824] do_msgsnd: ffff88001d9c9b80
[ 6.057224] do_msgsnd: ffff88001d9c9980
[ 6.057641] do_msgsnd: ffff88001d9c9c80
[ 6.058039] do_msgsnd: ffff88001d9c9600
[ 6.058415] do_msgsnd: ffff88001d9c9880
[ 6.058821] do_msgsnd: ffff88001d9c9700
[ 6.059528] do_msgsnd: ffff88001d9c9e80
[ 6.060729] do_msgsnd: ffff88001d6a4000
[ 6.061977] do_msgsnd: ffff88001d6a4080
[ 6.062912] do_msgsnd: ffff88001d6a4100
[ 6.063580] do_msgsnd: ffff88001d6a4180
[ 6.064034] do_msgsnd: ffff88001d6a4200
[ 6.064412] do_msgsnd: ffff88001d6a4280
[ 6.064889] do_msgsnd: ffff88001d6a4300
[ 6.065384] do_msgsnd: ffff88001d6a4380
[ 6.066055] do_msgsnd: ffff88001d6a4400
[ 6.066662] do_msgsnd: ffff88001d6a4480
[ 6.067127] do_msgsnd: ffff88001d6a4500
[ 6.067510] do_msgsnd: ffff88001d6a4580
[ 6.067929] do_msgsnd: ffff88001d6a4600
[ 6.068333] do_msgsnd: ffff88001d6a4680
```

- exhaust the *SLUB* freelist by allocating dozens of *msg_msg* objects
- this will make the heap layout more predictable
- and allows us to allocate adjacent objects

Freelist corruption - Exploit Section



- messages can be alloc'd via *msgsnd()*
- enqueued in a circular doubly linked list (*FIFO*)


Freelist corruption - Exploit Section



- they can be free'd specifically by requesting their *m_type* with *msgrcv()*
- for that, each message must have an unique *m_type*


Exploitation

Freelist corruption - Exploit Section




Exploitation

Freelist corruption - Exploit Section



Exploitation

Freelist corruption - Triggering the vulnerability



Exploitation

Freelist corruption - Triggering the vulnerability




Case 1

- *LSB* already zero
- vulnerable buffer free'd again
- nothing happens \Rightarrow retry!

Exploitation

Freelist corruption - Triggering the vulnerability




Case 2

- *LSB is 80*
- vulnerable buffer free'd again
- freelist pointer points to still alloc'd *msg_msg #3*

Exploitation

Freelist corruption - Double free



- regularly free *msg_msg #3*
- this causes a *double free* scenario
- and turns the freelist into a circular linked list
- similar to the *fastbin dup attack* against *glibc's ptmalloc2*

Freelist corruption - Double free

Great! We went from a poison nullbyte to a double free primitive. But what now?

The freelist holds two references to the same free'd object. We could

- alloc only one reference and tamper with freelist metadata (*UAF*)?
 - no leaks - directly writing freelist pointer not possible
 - we could overlap freelist pointer with object data
 - but I could not find a suitable object
- alloc both references and create a *type confusion*
 - overlap two different objects
 - sounds better, so I went with that

Type confusion

At this point, my goal was to leak some kernel pointers. And I had a vague plan

- allocate *type A*
 - *type A* must be somehow readable by userspace. Either directly or through a side channel
 - *type A* must not be free'd again but persist
- allocate *type B* on top of *type A*
 - *type B* must contain kernel pointers
- read *type A*



Type confusion - Candidates

- took a long time to find suitable candidates
- in the end, I identified two potential candidates via *pahole*

Type confusion - Type A

ip6_sf_socklist as *type A* candidate

- can be alloc'd via *setsockopt*
- content can be read back via *getsockopt*
- can be alloc'd in *kmalloc-128*
- unprivileged access

And as a bonus

- from *ip6_sf_socklist+0x8* content fully user controlled
- variable length, so it can be alloc'd in other caches, too

Type confusion - Type A

```
struct ip6_sf_socklist
    +0x0000 sl_max           : unsigned int
    +0x0004 sl_count        : unsigned int
    +0x0008 sl_addr         : struct in6_addr []
```

Type confusion - Type B

rfkill_data as *type B* candidate

- alloc'd by opening */dev/rfkill*
- contains multiple heap pointers
 - *rfkill_data.events*
 - *rfkill_data.mtx.wait_list*
 - *rfkill_data.read_wait.task_list*
- contains kernel text leak
 - *rfkill_data.list.prev* (*rfkill_fds*)
- most likely unprivileged access

There are some restrictions though

- *CONFIG_RFKILL=y* - no text leak if compiled as module
- */dev/rfkill* must be at least openable by unprivileged users

Type confusion - Type B

```
struct rfkill_data
+0x0000 list                : struct list_head
+0x0010 events              : struct list_head
+0x0020 mtx                 : struct mutex
+0x0048 read_wait           : wait_queue_head_t
+0x0060 input_handler       : bool
```

Exploitation

Type confusion - Type A and Type B

```
struct ip6_sf_socklist
+0x0000 sl_max           : unsigned int
+0x0004 sl_count         : unsigned int
+0x0008 sl_addr          : struct in6_addr []
```

```
struct in6_addr {
    union {
        __u8           u6_addr8[16];
#ifdef __UAPI_DEF_INET6_ADDR_ALT
        __be16         u6_addr16[8];
        __be32         u6_addr32[4];
#endif
    } in6_u;
};
```

```
struct rfkill_data
+0x0000 list             : struct list_head
+0x0010 events           : struct list_head
+0x0020 mtx              : struct mutex
+0x0048 read_wait        : wait_queue_head_t
+0x0060 input_handler     : bool
```

```
struct list_head
+0x0000 next             : struct list_head *
+0x0008 prev             : struct list_head *
```

- `.sl_addr` readable via `getsockopt`
- *16-bytes* per entry

⇒ we can read `rfkill_data` starting from offset `0x8`

Exploitation

Type confusion - Kernel text leak

```
struct ip6_sf_socklist
+0x0000 sl_max           : unsigned int
+0x0004 sl_count         : unsigned int
+0x0008 sl_addr          : struct in6_addr []
```

```
struct in6_addr {
    union {
        __u8           u6_addr8[16];
#ifdef __UAPI_DEF_INET6_ADDR_ALT
        __be16         u6_addr16[8];
        __be32         u6_addr32[4];
#endif
    } in6_u;
};
```

```
struct rfkill_data
+0x0000 list             : struct list_head
+0x0010 events           : struct list_head
+0x0020 mtx              : struct mutex
+0x0048 read_wait        : wait_queue_head_t
+0x0060 input_handler    : bool
```

```
struct list_head
+0x0000 next             : struct list_head *
+0x0008 prev            : struct list_head *
```

- `.sl_addr[0] = rfkill_data.list.prev`
- upon creation, always points to `rfkill_fds`
- kernel text base can be derived from `rfkill_fds`

Exploitation

Type confusion - kcalloc-128 leak

```
struct ip6_sf_socklist
+0x0000 sl_max          : unsigned int
+0x0004 sl_count        : unsigned int
+0x0008 sl_addr         : struct in6_addr []
```

```
struct in6_addr {
    union {
        __u8          u6_addr8[16];
#ifdef __UAPI_DEF_INET6_ADDR_ALT
        __be16        u6_addr16[8];
        __be32        u6_addr32[4];
#endif
    } in6_u;
};
```

```
struct rfkill_data
+0x0000 list            : struct list_head
+0x0010 events          : struct list_head
+0x0020 mtx             : struct mutex
+0x0048 read_wait       : wait_queue_head_t
+0x0060 input_handler   : bool
```

```
struct list_head
+0x0000 next            : struct list_head *
+0x0008 prev           : struct list_head *
```

- multiple ways to leak *kmalloc-128* base via empty lists
 - *rfkill_data.events* (might not be empty)
 - *rfkill_data.mtx.wait_list* (most likely empty)
 - *rfkill_data.read_wait.task_list* (probably empty)

Escalating privileges

Now that we have leaked some kernel pointers, let's try to escalate to root!

Escalating privileges - Idea #1

Idea: Directly set *task_struct.cred* pointer to root creds

- global data pointer to *init_cred* already leaked
- requires arbitrary write primitive
- requires *task_struct* base address
 - maybe through *rfkill_data.mtx.owner*
 - thread A acquires *mutex*, thread B leaks *.owner* via type confusion
 - very tiny window - unlikely to win the race
- decided against this idea

Escalating privileges - Idea #2

Idea: Hijack *RIP* and go for classic *prepare_kernel_cred/commit_creds*

- kernel text leaked, so we could *ROP*
- requires write against a function pointer
- sounds promising enough

RIP hijack

- using *type confusion*, it is possible to overwrite *kmalloc-128* objects
 - alloc *ip6_sf_socklist* on top of target object
 - remember: *ip6_sf_socklist.sl_addr* is user controlled
- but, could not find target object in *kmalloc-128* with function pointers
- instead, maybe target global function pointers?
 - *ptmx_fops* could work
 - easily triggered by performing file ops on */dev/pmtx*
 - but requires arbitrary write

Exploitation

RIP hijack - Arbitrary write

Took a while, but I came up with a plan on how to craft an arbitrary write.

RIP hijack - Arbitrary write

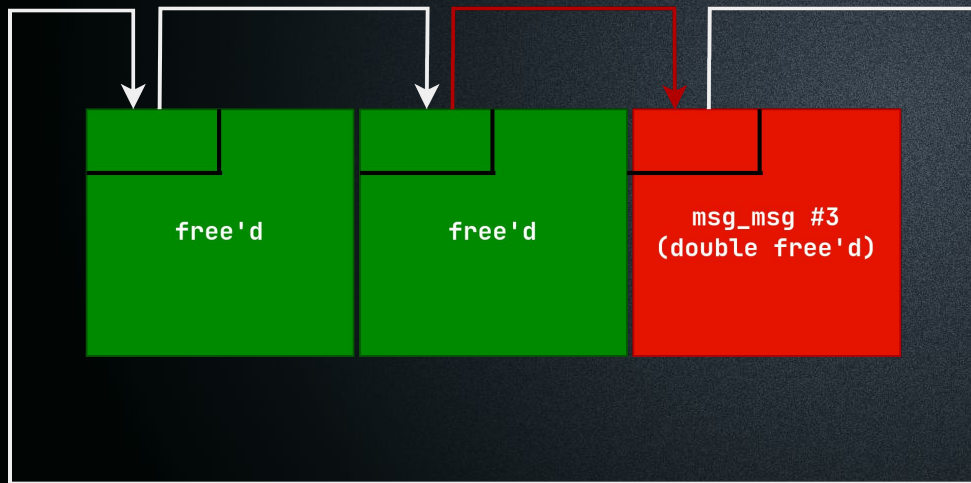
- corrupt a freelist pointer to enable arbitrary allocations
- allocate an *ip6_sf_socklist* object on top of *ptmx_fops*
 - *ip6_sf_socklist* holds user controlled data at offset *0x8*

RIP hijack - Arbitrary write

But how can we overwrite a freelist pointer?

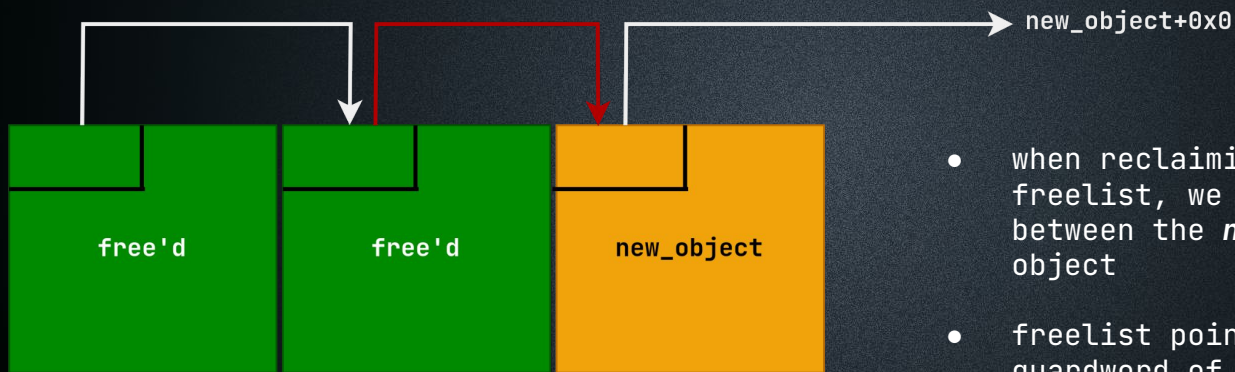
Exploitation

RIP hijack - Arbitrary write



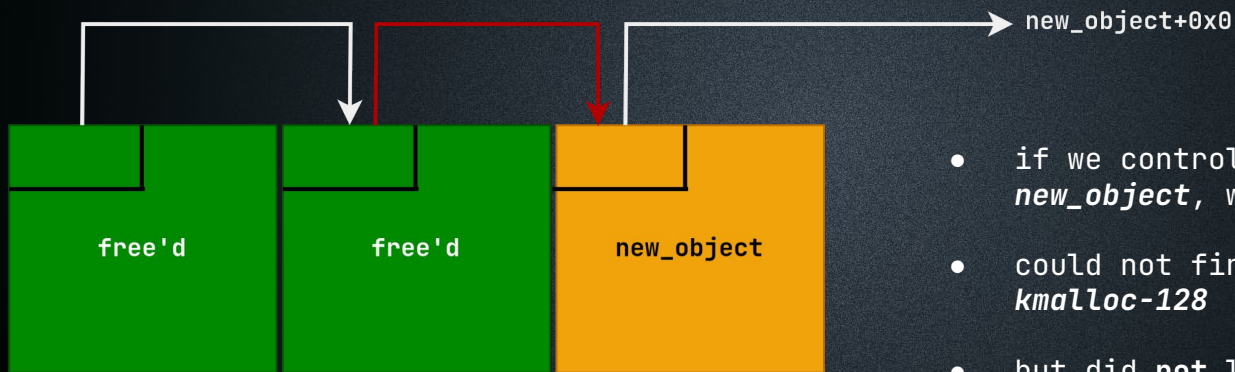
- repeat the double free

RIP hijack - Arbitrary write



- when reclaiming *msg_msg* #3 from the freelist, we create a *type confusion* between the *new_object* and a *free'd* object
- freelist pointer will be set to first quardword of the *new_object*

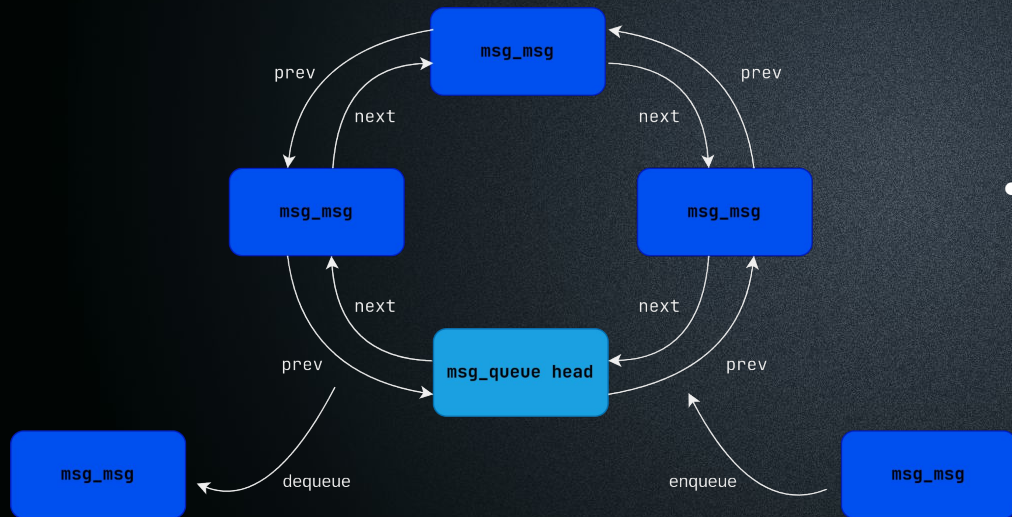
RIP hijack - Arbitrary write



- if we control the first quadword of *new_object*, we control the freelist
- could not find a suitable object in *kmalloc-128*
- but did **not** look into other caches

Exploitation


RIP hijack - Arbitrary write



- instead, abuse `msg_queue` linking process

Exploitation


RIP hijack - Arbitrary write



```
struct msg_msg
+0x0000 m_list      : struct list_head
+0x0010 m_type      : long
+0x0018 m_ts        : size_t
+0x0020 next        : struct msg_msgseg *
+0x0028 security    : void *
```

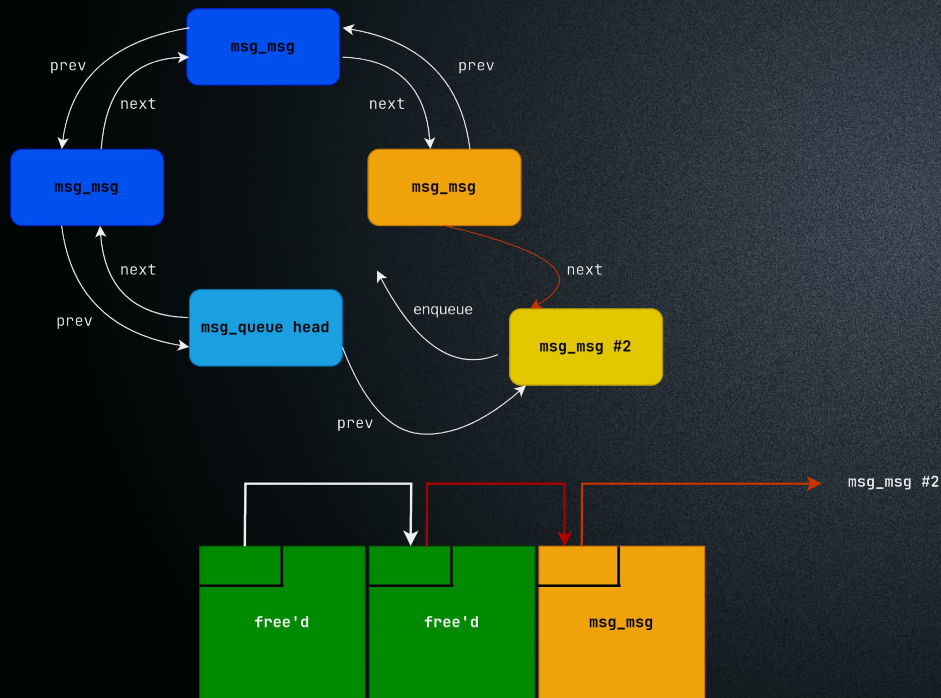
```
struct list_head
+0x0000 next        : struct list_head *
+0x0008 prev        : struct list_head *
```

- choose *msg_msg* as *new_object*
- freelist pointer repurposed as *.m_list.next* pointer
- freelist merged with *msg_queue*



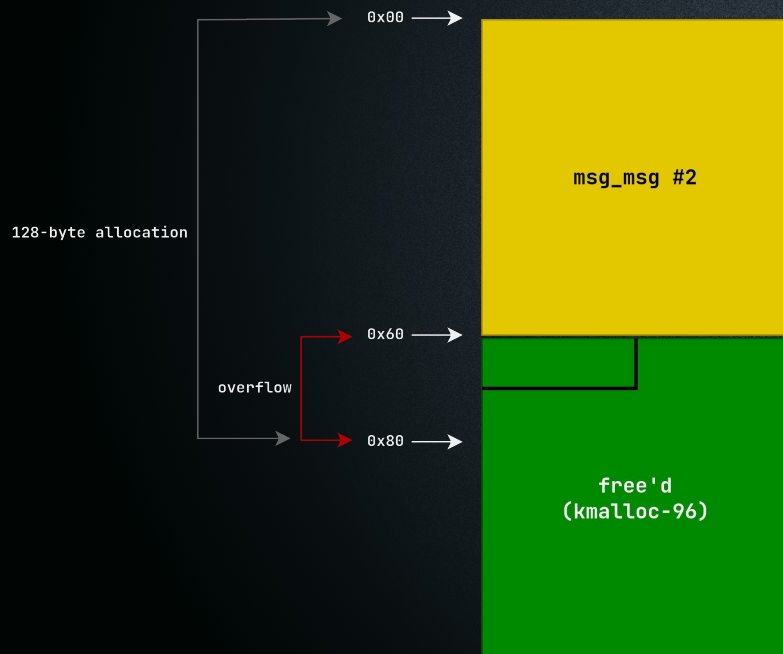
Exploitation

RIP hijack - Arbitrary write



- `msg_msg #2` should live in `< kmalloc-128`
 - in our case: `kmalloc-96`
- `96-bytes` chunk linked in `128-bytes` freelist

RIP hijack - Arbitrary write



- size mismatch when requesting **128 bytes** from the allocator
- giving us an overflow of **32 bytes**
- enough to overwrite **kmalloc-96** freelist pointer

RIP hijack - Arbitrary write

- now we can perform (nearly) arbitrary *96-bytes* allocations
- combining that with *ip6_sf_socklist*, we have our arbitrary write!

RIP hijack - Arbitrary write

Before we resume, I would like to highlight the drawbacks of this approach...

- rather invasive
- leaves some kernel structures in a weird state
- that means more cleanup afterwards
- approach avoidable, if we find a kernel object with user controlled first quadword
 - that would also reduce exploit complexity

RIP hijack

Using our arbitrary write primitive we can

- overwrite the first *96-bytes* of *ptmx_fops*
- and replace one of the callbacks with our *stack pivot* gadget

RIP hijack - Stack pivot

- luckily I found following gadget

```
push rdx
mov edx, 0x415bffb7
pop rsp
pop rbp
ret
```

- overwrite *ptmx_fops.unlocked_ioctl* with stack pivot gadget
 - it lets us pass *8-bytes* of user data to the gadget via *rdx*
 - perfect to smuggle a *kernel space* pointer

```
SYSCALL_DEFINE3(ioctl, unsigned int, fd, unsigned int, cmd, unsigned long, arg)
```


RIP hijack - Stack pivot

- choose address of new fake stack
- the previously leaked *kmalloc-128* pointer is a good choice
 - we can setup the stack by replacing currently alloc'd object with *ip6_sf_socklist*

All we have to do now, is to *ioctl* on */dev/ptmx*

```
ioctl(tty_fd, 0xbeef, kptr_kmalloc128 + 0x8);
```

RIP hijack - Stage-1 ROP chain

- we let *rsp* point to a *ip6_sf_socklist* object in *kmalloc-128* where we have already prepared our first stage rop chain
- the new fake stack might be end of page
 - meaning, we can safely use **120** bytes for the rop chain
 - or in gadgets: **15** gadgets
- that's not a lot...

RIP hijack - Stage-1 ROP chain

- since we have only **15** gadgets, we will pivot the stack one more time
- this time we pivot to the kernel log buffer (***__log_buf***)
 - within the kernel **BSS**
 - safe to repurpose as stack
 - of course not the stealthiest way (***dmesg*** truncated)
- before pivot, we copy the second stage from userspace to ***__log_buf***

RIP hijack - Stage-1 ROP chain

```
<rbp = dummy>
<call_rwsem_down_read_failed+34> → pop    rsi
                                   pop    rdi
                                   pop    rbp
                                   ret

<rsi = &rop_chain_second_stage>
<rdi = __log_buf+2048>
<rbp = dummy>
<intel_edp_backlight_power+69> → pop rdx
                                   ret

<rdx = sizeof(rop_chain_second_stage)>
<_copy_from_user>
<bstat+52> → leave
               ret
```


RIP hijack - Stage-2 ROP chain

The second stage takes care of

- restore *ptmx_fops*
- remove allocated *ip6_sf_socklist* objects
- attach *root creds* to current task
- exit *kernel space* (*swapgs, iretq*)

Demo

DEMO TIME

Spawning a root shell

- high reliability
- roughly 94.44% chance of success

Stability afterwards

- 10% to 20% chance of crashing the kernel
- improvable by further cleanup

Final words

Thanks!

Exploit: https://github.com/Milo-D/CVE-2016-6187_LPE